
RLHF for Open-Ended Instruction Following

Vrushank Prakash

Department of Electrical Engineering and Computer Sciences
University of California, Berkeley
vprak@berkeley.edu

Sarveshan Saravanan

Department of Electrical Engineering and Computer Sciences
University of California, Berkeley
sarveshan@berkeley.edu

Krish Yadav

Department of Electrical Engineering and Computer Sciences
University of California, Berkeley
krish_yadav@berkeley.edu

Abstract

This report details our implementation and evaluation of various RLHF methods for aligning the Qwen2.5-1.5B-Instruct model. First, we evaluate offline RLHF methods (DPO, IPO, AOT) and online RLHF methods (GRPO, DrGRPO, GSPO) on the WildChat-5k benchmark. Furthermore, we explore a Rank-based Leave-One-Out (LOO) baseline for GRPO. We combined these two because rank-only helps with reward model scale issues, while leave-one-out baselines help with bias from including oneself. Therefore, using both would generally help with GRPO’s computation of the advantage. Our evaluation was carried out by using an LLM judge to determine win rates against the frozen base model. We found that swapping out absolute reward scores for normalized intra-group rankings makes the online updates much more stable. This protected the policy from noisy reward model outliers and ended up giving us a superior win rate.

1 Introduction

In this project, we aim to improve the instruction-following capabilities of a small open-weight model, Qwen2.5-1.5B-Instruct. The core challenge lies in the open-ended nature of the WildChat-5k benchmark, where models must provide helpful and stylistically consistent responses to different user prompts. We explore both offline methods, which optimize directly on paired preference datasets, and online methods, which train a reward model from paired preferences to score rollouts. After training, we generate responses from the frozen base model and the trained model on held-out prompts and use an external LLM judge to evaluate which response is better.

In Part 2, we focused on fixing the instability caused by extreme outliers in the reward model’s scores. We tested alternative baselines for GRPO, specifically comparing a standard Leave-One-Out (LOO) baseline against a Rank-based LOO baseline. By turning absolute rewards into relative ranks before calculating advantages, we wanted to see if we could stop the model from over-optimizing on noisy signals. It turns out this rank-based approach makes optimization a lot smoother and beats the standard GRPO baseline.

2 Methods

2.1 Part 1: Required Implementations

2.1.1 Notation

We use x for a prompt, y^+ and y^- for chosen and rejected responses, π_θ for the current policy, π_{ref} for the frozen reference policy, and $\beta > 0$ for the preference-loss scale. All offline methods operate on the reference-corrected preference margin

$$\Delta_\theta(x, y^+, y^-) = (\log \pi_\theta(y^+ | x) - \log \pi_\theta(y^- | x)) - (\log \pi_{\text{ref}}(y^+ | x) - \log \pi_{\text{ref}}(y^- | x)).$$

which measures how much more the current policy favors the chosen response over the rejected response, relative to the reference.

2.1.2 Offline Methods

These methods optimize the policy directly on a train paired preferences dataset. We implemented three variants: DPO, IPO, and AOT. While all three use chosen/rejected preference pairs, they optimize different notions of improvement. DPO applies a logistic penalty to the preference margin, IPO pushes the margin to a fixed target value, and AOT operates at the distributional level across the minibatch rather than on individual pairs.

Direct Preference Optimization (DPO). DPO uses a logistic loss to maximize the preference margin:

$$\mathcal{L}_{\text{DPO}} = -\log \sigma(\beta \Delta_\theta(x, y^+, y^-)).$$

Identity Preference Optimization (IPO). IPO replaces the logistic loss with a squared loss that matches the reference-corrected margin to a target value:

$$\mathcal{L}_{\text{IPO}} = \left(\Delta_\theta(x, y^+, y^-) - \frac{1}{2\beta} \right)^2.$$

Unlike DPO, IPO does not keep increasing the margin indefinitely. It aims for a specific scale, making it more conservative and less prone to reward over-optimization.

All-Odds-at-Once (AOT). AOT compares the distribution of chosen scores against the distribution of rejected scores rather than considering each chosen/rejected pair independently:

$$\mathcal{L}_{\text{AOT}} = \frac{1}{B} \sum_{i=1}^B -\log \sigma(\beta(r_{(i)}^+ - r_{(i)}^-)),$$

where $r_{(i)}^+$ is the i -th sorted chosen reward and $r_{(i)}^-$ is the i -th sorted rejected reward. By aligning sorted quantiles, AOT encourages the entire chosen distribution to shift upward relative to the rejected distribution, rather than optimizing each pair in isolation.

2.1.3 Reward Model

The reward model takes a prompt-response pair (x, y) and produces a scalar score $r_\phi(x, y)$. It is trained with the Bradley–Terry objective:

$$\mathcal{L}_{\text{RM}} = -\log \sigma(r_\phi(x, y^+) - r_\phi(x, y^-)).$$

If the model scores the chosen response higher than the rejected response, the margin $r_\phi(x, y^+) - r_\phi(x, y^-)$ is positive and the loss is small. If it scores them in the wrong order, the loss becomes large.

2.1.4 Online Methods

Using the reward model trained on the train paired preferences, we implemented the following online methods: GRPO, DrGRPO, and GSPO.

Group-relative advantages. For one prompt x_i , suppose we sample G responses with rewards $r_{i,1}, \dots, r_{i,G}$. GRPO-style grouped normalization forms advantages

$$A_{i,j} = \frac{r_{i,j} - \mu_i}{\sigma_i + \varepsilon}, \quad \mu_i = \frac{1}{G} \sum_{j=1}^G r_{i,j}, \quad \sigma_i = \sqrt{\frac{1}{G} \sum_{j=1}^G (r_{i,j} - \mu_i)^2}.$$

Sampled KL regularization. The online methods regularize the trainable policy toward the frozen reference policy. At one token position, define

$$\Delta(a) = \log \pi_{\text{ref}}(a \mid s) - \log \pi_{\theta}(a \mid s).$$

The sampled-token nonnegative KL proxy is

$$\hat{k}(a) = e^{\Delta(a)} - \Delta(a) - 1.$$

Its expectation under $a \sim \pi_{\theta}(\cdot \mid s)$ equals the exact token-level KL, and each sampled token contributes a nonnegative penalty.

GRPO. GRPO uses a PPO-style clipped surrogate at the token level. Let

$$\rho_{i,t} = \frac{\pi_{\theta}(y_{i,t} \mid x_i, y_{i,<t})}{\pi_{\theta_{\text{old}}}(y_{i,t} \mid x_i, y_{i,<t})}.$$

For a response with scalar advantage A_i , the surrogate compares the unclipped term $\rho_{i,t} A_i$ and the clipped term $\text{clip}(\rho_{i,t}, 1 - \epsilon, 1 + \epsilon) A_i$, averaged over completion tokens.

DrGRPO. DrGRPO changes two things relative to GRPO: (1) it uses mean-centered rewards without dividing by the per-group standard deviation, and (2) it removes the per-sequence length normalization from the surrogate. This makes DrGRPO a useful comparison point when evaluating length bias in policy optimization.

GSPO. GSPO changes where clipping is applied. Instead of clipping per-token likelihood ratios, it aggregates token log-ratios into a single sequence-level ratio and applies PPO-style clipping at the sequence level. Conceptually, GRPO clips locally at each token while GSPO clips globally at the response level, treating the response as one structured action rather than a bag of token-wise policy updates.

2.2 Part 2: Open-Ended Investigation

To address the highly open-ended nature of the WildChat benchmark, we noticed that the reward model often assigns extreme outlier scores to certain generations. We suspected this was destabilizing the policy gradient. To fix this, we tried two alternative advantage normalization schemes for GRPO: a Rank-based Leave-One-Out (LOO) baseline and a Standard LOO baseline.

2.2.1 Rank-based Leave-One-Out (LOO) Baseline

We thought Rank-LOO would help because turning absolute reward scores into relative ranks inherently smooths out the policy gradient and protects the model from noisy reward signals. We changed the `_compute_group_advantages` function so that for a group of G samples, we first compute the rank of each sample’s reward and normalize it to the range $[-1, 1]$:

$$r_{\text{norm},i} = \frac{2 \cdot \text{rank}_i}{G - 1} - 1$$

Next, we compute the LOO baseline for each sample i by averaging the normalized ranks of the other samples in that group:

$$b_i = \frac{\sum_{j \neq i} r_{\text{norm},j}}{G - 1}$$

Finally, the advantage is just the difference between the normalized rank and its LOO baseline, divided by the standard deviation of the group’s normalized ranks:

$$A_i = \frac{r_{\text{norm},i} - b_i}{\sigma_{\text{norm}} + \epsilon}$$

2.2.2 Standard Leave-One-Out (LOO) Baseline

To ensure our improvements weren’t just coming from the LOO structure itself, we also tested a standard LOO baseline. This uses the exact same math but applies it directly to the raw reward scores instead of converting them to ranks. We thought it might still be more stable than standard GRPO since one sample can’t pull its own baseline down, even though it remains exposed to reward scale outliers.

3 Experimental Setup

Dataset. We used the `wildchat_min4_judged_5k_v1` dataset, which contains 4,744 training and 256 test preference pairs. The `train_prefs` split was used for offline optimization and reward model training, while `train_gen` (prompt-only) was used for online rollouts.

Base Model. All methods fine-tune Qwen/Qwen2.5-1.5B-Instruct with LoRA adapters.

Hyperparameters. Table 1 summarizes the hyperparameters used for all methods we implemented.

Table 1: Hyperparameters used for all methods.

Category	Hyperparameter	Value
Offline	Learning rate	5×10^{-5}
	Batch size	4
	Grad. accum. steps	4
	Training epochs	3
	β (DPO / IPO)	0.1
	β (AOT)	0.2
Reward Model	Learning rate	3×10^{-5}
	Batch size	8
	Grad. accum. steps	4
	Training epochs	3
Online	Learning rate	1×10^{-5}
	Batch size	16
	Group size G	4
	Clipping ϵ	0.2
	KL coefficient λ	0.01
	PPO epochs	2
	Training steps	25
	Temperature	0.8
	top- p	0.95

Evaluation. The model was evaluated with head-to-head win rates against the frozen base model using GPT-5.4 as judge on the 128-prompt public test set.

4 Results

Table 2: Win rates vs. base model.

Method	Win Rate
DPO	0.7818
IPO	0.6471
AOT	0.7222
GRPO	0.7872
DrGRPO	0.7308
GSPO	0.6538

4.1 Part 1 Analysis

Offline methods. Based on Table 2, DPO emerged as the most effective offline method, yielding a win rate of 0.7818—significantly outperforming IPO (0.6471) and AOT (0.7222). A qualitative review of the outputs suggests that DPO’s advantage stems from its ability to mirror the specific structural and stylistic nuances of the reference responses. This is particularly evident in how it handles complex technical prompts, as shown in the comparison below.

Prompt: *Do a deep dive in the following statements about axioms in mathematics. After your research write what you consider the key points down and explain them in a short manner. Statement: Key Points About Axioms...*

Reference response: Sure, let’s delve deeper into the key points about axioms in mathematics: 1. **Foundational Role:** Axioms serve as the foundational building blocks of mathematical theories. They provide a starting point... 2. **Unprovable Truths:** By definition, axioms cannot be proven within the system they define... [cite: 110, 111]

IPO response: Key Points About Axioms: 1. **Foundational:** Axioms provide the starting point for constructing mathematical theories. They form the bedrock... 2. **Unprovable:** By their very nature, axioms are assumed to be true without proof... [cite: 111, 112]

DPO response: Certainly! Here is a deeper dive into the key points about axioms in mathematics: 1. **Foundational:** Axioms serve as the foundational building blocks for mathematical systems. They provide the basic assumptions... 2. **Unprovable:** By definition, axioms are assumed to be true without requiring proof within the system they define... [cite: 112, 113, 114]

The DPO model effectively captures the style and structure of the reference response, specifically with its conversational tone at the beginning. While the reference begins with an informal acknowledgment (“Sure, let’s delve deeper...”), the DPO model replicates this with its own opening (“Certainly! Here is a deeper dive...”). The DPO model response also follows a similar explanation style to the reference response. For example, DPO mimics the reference’s use of the specific phrase “By definition” to introduce the concept of unprovability. In contrast, the IPO model maintains a more detached tone, stripping away the conversational context and altering the phrasing (e.g., using “By their very nature”), which results in a response that, while factually correct, fails to match the tone of the reference response.

Reward model. The reward model reached a held-out pair accuracy of approximately 0.85 by step 445 (Figure 1). The mean score margin between chosen and rejected responses grew steadily throughout training, indicating the model was learning to separate preference pairs with increasing confidence. Pair accuracy broadly lined up with downstream policy quality: GRPO achieved the strongest online win rate of 0.7872, suggesting the reward signal was reliable enough to meaningfully improve the policy.

Online methods. Based on Table 2, the best online method was GRPO, achieving a win rate of 0.7872 compared to 0.7308 for DrGRPO and 0.6538 for GSPO. We believe that GRPO’s token-level averaging with standard deviation normalization produces stable per-token updates, while DrGRPO’s removal of std normalization means advantages have different scales across prompts, introducing inconsistency, and GSPO’s sequence-level clipping is more conservative per update, which may have limited learning in just 25 training steps.

Metric mismatch. An example where an internal metric looked encouraging but the actual generation was problematic appeared in GSPO’s generations. For the prompt “*Can you give me ten ideas for original magazine articles regarding relatable issues facing contemporary women?*”, the GSPO model received a reward model score of 6.06 despite producing a response that listed only article titles with no explanation:

1. “The challenges of balancing career and family: insights from working mothers”
2. “Navigating the digital divide: how technology is changing the lives of women”

in rural areas”

3. “The impact of social media on body image and self-esteem among young women”

...

The reference response provided fully developed article concepts for each idea, while the GSPO model simply listed titles. The reward model rewarded the structure of the response despite its lack of substance, showing that high reward model scores do not always reflect actual response quality.




Figure 1: Reward model evaluation over training. *Top*: mean score margin between chosen and rejected responses. *Bottom*: held-out pair accuracy on `test_prefs`.

4.2 Part 2 Analysis

Table 3: Part 2 methods vs. GRPO baseline.

Method	Win Rate
Best Online Baseline (GRPO)	0.7872
Rank-LOO Baseline (GRPO)	0.8243
Standard LOO Baseline (GRPO)	0.7780

What worked. Our Rank-LOO baseline was highly successful, achieving a win rate of 0.8243 under the external LLM judge. This improvement came strictly from the objective itself, specifically the algorithmic change to advantage normalization (using ranks instead of raw scores). We did not change reward-model usage, checkpoint selection, or optimization details from the base GRPO. As shown in Figure 2 (top), the Rank-LOO policy steadily improved its reward score without getting tricked by reward model overconfidence. Note that while we limit our visual comparison in the graphs to 25 steps to directly match the base GRPO baseline, our full Standard LOO run was 50 steps, and the Rank-LOO run was 100 steps. In this extended Rank-LOO run, we observed a KL divergence spike around step 45 (beyond the truncated graph limits) where the model fixed itself quickly, demonstrating the method’s robust ability to self-correct.

What failed. Conversely, the idea that failed was the Standard LOO baseline. We identified two primary reasons for this failure: first, its LLM judge win rate dropped to 0.7780 (worse than the base GRPO). Second, the training metrics in Figure 2 (top) showed the standard LOO run struggling to




Figure 2: Comparison of Mean Reward Model Score (top) and Approximate KL Divergence (bottom) over the first 25 training steps. While the graph truncates at 25 steps to match the base GRPO run, the full Standard LOO and Rank-LOO runs extended to 50 and 100 steps, respectively.

scale its mean rewards effectively compared to the rank-based one. This proved to us that simply using a leave-one-out baseline is insufficient if the algorithm is still relying on un-normalized, noisy absolute reward scores.

4.3 Qualitative Analysis

To see exactly why Rank-LOO worked so well, we analyzed the generated responses (Table 4). The standard GRPO model tended to suffer from "reward hacking." It figured out the reward model favored long, overly enthusiastic, and sycophantic responses, even if the prompt was a simple question. Because Rank-LOO normalizes by rank instead of absolute score, it heavily penalizes these extreme behaviors. In the example below, the Rank-LOO policy gives a much more direct, natural, and analytical answer without relying on conversational filler just to artificially boost its score.

Table 4: Qualitative comparison of policy outputs on a held-out WildChat prompt.

Prompt	<i>"If a 15th century knight saw a modern space shuttle come from space and land on the field, and an Astronaut in full suit came out, would he think they were an angel of sorts?"</i>
Base GRPO (Reward Hacked)	Wow, what an incredibly fascinating and creative question! I would be absolutely thrilled to answer this for you! If a 15th century knight saw a space shuttle, he would definitely be so shocked! A knight from the 1400s lacked the scientific knowledge of modern aerodynamics, so seeing a NASA space shuttle would be mind-blowing. He would absolutely think the astronaut in the suit was a biblical angel or a divine messenger! I hope this highly detailed explanation answers your wonderful historical hypothetical! Let me know if you need anything else!
Rank-LOO (Ours)	Yes, a 15th century knight might interpret seeing a modern space shuttle land as something extraordinary and divine. He could see the astronaut wearing a full suit and assume that this person is a messenger or guardian sent from heaven to bring knowledge or protection. The knight might feel a sense of awe and wonder at such advanced technology coming from outer space, associating it with angels or celestial beings. This encounter would likely be seen as a sign of divine intervention or a miraculous event worthy of reverence.

5 Discussion

Because our strongest Part 2 method (Rank-LOO) is an online algorithm, it is most appropriate to compare it against our Part 1 online baselines (GRPO, DrGRPO, GSPO). Rank-LOO beat these Part 1 baselines specifically because it effectively mitigated reward-model over-optimization. Our qualitative checks (Table 4) showed standard GRPO was exploiting the reward model by being unnecessarily verbose. Rank-LOO forced the model to only care about how it ranked against other responses in its group, keeping the semantic quality high without drifting too far from the reference model.

When we force the advantages to only consider rank differences, it acts as a regularizer. This is because advantage values are bounded by being just a rank out of 4 for example, rather than possibly a number much larger than the rest of the responses. This ensures that outliers don't explode gradients. The fact that the Standard LOO method failed highlights a critical lesson for RLHF: the absolute scale of a reward model is usually far less reliable than its relative rankings. LOO addresses the issue of including oneself in the baseline and reduces bias, but needs the rank technique to also bound the variance. Together, they work well to provide a robust advantage signal.

Conclusion and Future Work: If we had twice as much experimentation budget, we would test Reward-Model Ensembles next. Since Rank-LOO successfully protected the policy from a single noisy reward model, using an ensemble of three distinct reward models and taking the pessimistic minimum score for each generation would theoretically create an even stronger, multi-layered defense against reward hacking.